

Stanford CS224W: Recommender Systems: Task and Evaluation

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>




Preliminary of Recommendation

- **Information Explosion in the era of Internet**
 - 10K+ movies in Netflix
 - 12M products in Amazon
 - 70M+ music tracks in Spotify
 - 10B+ videos on YouTube
 - 200B+ pins (images) in Pinterest
- **Personalized recommendation (i.e., suggesting a small number of interesting items for each user)** is critical for users to effectively explore the content of their interest.

Recommender System as a Graph

- Recommender system can be naturally modeled as a **bipartite graph**
 - A graph with two node types: **users** and **items**.
 - **Edges** connect users and items
 - Indicates user-item interaction (e.g., click, purchase, review etc.)
 - Often associated with timestamp (timing of the interaction).




Recommendation Task

■ Given

- Past user-item interactions

■ Task

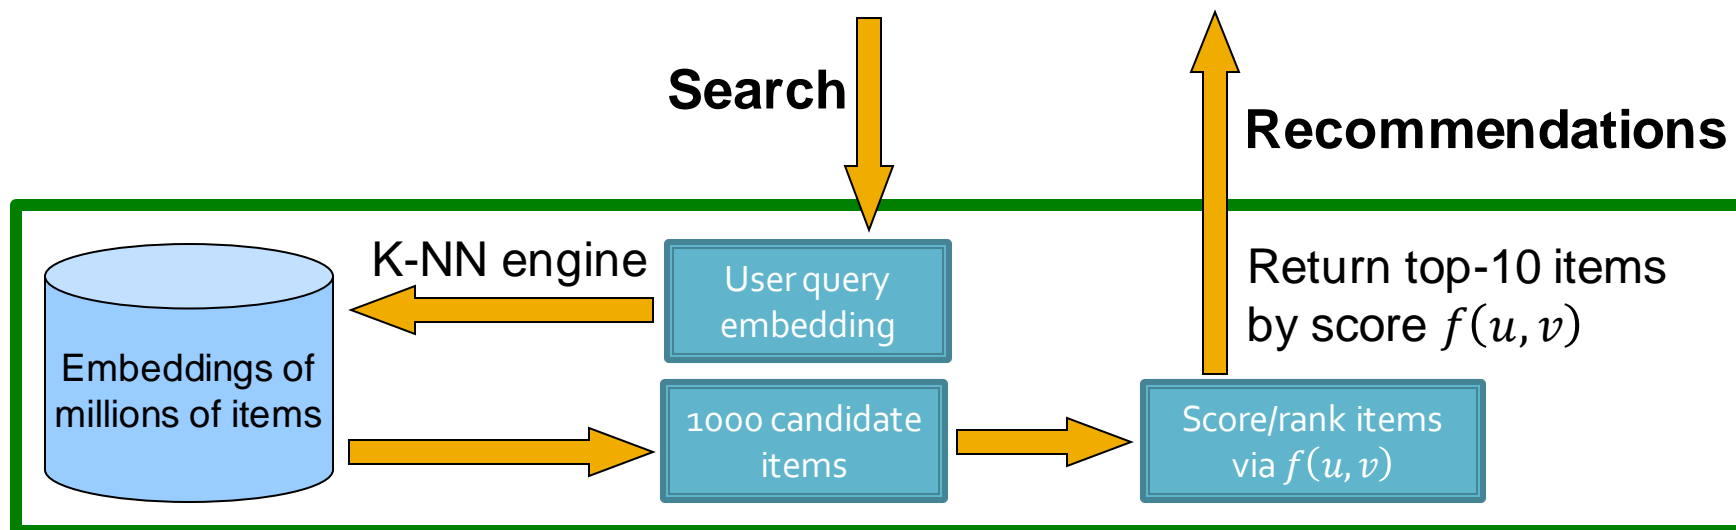
- Predict new items each user will interact in the future.
- Can be cast as **link prediction** problem.
 - Predict new user-item interaction edges given the past edges.
- For $u \in U, v \in V$, we need to get a real-valued **score** $f(u, v)$.



Modern Recommender System

- **Problem:** Cannot evaluate $f(u, v)$ for every user u – item v pair.
- **Solution:** 2-stage process:
 - Candidate generation (cheap, fast)
 - Ranking (slow, accurate)

Example $f(u, v)$:
 $f(u, v) = z_u \cdot z_v$




Top- K Recommendation

- For each user, we recommend K items.
 - For recommendation to be effective, K needs to be much smaller than the total number of items (up to billions)
 - K is typically in the order of 10—100.
- The goal is to include as many **positive items** as possible in the top- K recommended items.
 - **Positive items = Items that the user will interact with in the future.**
- **Evaluation metric:** Recall@ K (defined next)


Evaluation Metric: Recall@K (1)

- For each user u ,
 - Let P_u be a set of positive items the user will interact in the future.
 - Let R_u be a set of items recommended by the model.
 - In top- K recommendation, $|R_u| = K$.
 - Items that the user has already interacted are excluded.



Evaluation Metric: Recall@K (2)

- **Recall@K** for user u is $|P_u \cap R_u| / |P_u|$.
 - Higher value indicates more positive items are recommended in top- K for user u .



- The final Recall@K is computed by averaging the recall values across all users.

Stanford CS224W: Recommender Systems: Embedding-Based Models

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>




Notation

■ Notation:

- U : A set of all users
- V : A set of all items
- E : A set of observed user-item interactions
 - $E = \{(u, v) \mid u \in U, v \in V, u \text{ interacted with } v\}$

Score Function


- To get the top- K items, we need a score function for user-item interaction:
 - For $u \in U$, $v \in V$, we need to get a real-valued scalar $\text{score}(u, v)$.
 - **K items with the largest scores for a given user u** (excluding **already-interacted items**) are then recommended.



For $K = 2$, recommended items for user u would be $\{v_1, v_3\}$.

Embedding-Based Models

- We consider **embedding-based models** for scoring user-item interactions.
 - For each user $u \in U$, let $\mathbf{u} \in \mathbb{R}^D$ be its D -dimensional embedding.
 - For each item $v \in V$, let $\mathbf{v} \in \mathbb{R}^D$ be its D -dimensional embedding.
 - Let $f_\theta(\cdot, \cdot): \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}$ be a parametrized function.
 - Then, $\text{score}(u, v) \equiv f_\theta(\mathbf{u}, \mathbf{v})$



Training Objective


- Embedding-based models have three kinds of parameters:
 - An encoder to generate user embeddings $\{\mathbf{u}\}_{u \in U}$
 - An encoder to generate item embeddings $\{\mathbf{v}\}_{v \in V}$
 - Score function $f_{\theta}(\cdot, \cdot)$
- **Training objective**: Optimize the model parameters to achieve high recall@ K on *seen* (i.e., *training*) user-item interactions
 - We hope this objective would lead to high recall@ K on *unseen* (i.e., *test*) interactions.

Surrogate Loss Functions

- The original training objective ($\text{recall}@K$) is **not differentiable**.
 - *Cannot* apply efficient gradient-based optimization.
- Two **surrogate loss functions** are widely-used to enable efficient gradient-based optimization.
 - Binary loss
 - Bayesian Personalized Ranking (BPR) loss
- Surrogate losses are **differentiable** and should **align well with the original training objective**.

Binary Loss (1)

- Define **positive/negative edges**
 - A set of **positive edges** E (i.e., observed/training user-item interactions)
 - A set of **negative edges** $E_{\text{neg}} = \{(u, v) \mid (u, v) \notin E, u \in U, v \in V\}$
- Define **sigmoid function** $\sigma(x) \equiv \frac{1}{1+\exp(-x)}$
 - Maps real-valued scores into binary likelihood scores, i.e., in the range of $[0,1]$.



Binary Loss (2)

- **Binary loss**: Binary classification of **positive/negative** edges using $\sigma(f_\theta(\mathbf{u}, \mathbf{v}))$:

$$\underbrace{-\frac{1}{|\mathbf{E}|} \sum_{(u,v) \in \mathbf{E}} \log(\sigma(f_\theta(\mathbf{u}, \mathbf{v})))}_{\text{positive edges}} - \underbrace{\frac{1}{|\mathbf{E}_{\text{neg}}|} \sum_{(u,v) \in \mathbf{E}_{\text{neg}}} \log(1 - \sigma(f_\theta(\mathbf{u}, \mathbf{v})))}_{\text{negative edges}}$$

During training, these terms can be approximated using mini-batch of positive/negative edges


- Binary loss pushes the scores of **positive edges** higher than those of **negative edges**.
 - This aligns with the training recall metric since positive edges need to be recalled.

Issue with Binary Loss (1)

- **Issue:** In the binary loss, the scores of **ALL** positive edges are pushed higher than those of **ALL** negative edges.
- This would unnecessarily penalize model predictions even if the training recall metric is perfect.
- **Why?** (example in the next slide)

Issue with Binary Loss (2)

- Let's consider the simplest case:
 - Two users, two items
 - Metric: Recall@1.
 - A model assigns the score for every user-item pair (as shown in the right).
- Training **Recall@1 is 1.0** (perfect score), because v_0 (resp. v_1) is correctly recommended to u_0 (resp. u_1).
- However, **the binary loss would still penalize the model prediction** because the negative (u_1, v_0) edge gets the higher score than the positive edge (u_0, v_0) .



Issue with Binary Loss (3)


- **Key insight:** The binary loss is **non-personalized** in the sense that the **positive/negative edges are considered *across ALL users at once***.
- However, the recall metric is inherently **personalized (defined for each user)**.
 - The non-personalized binary loss is overly-stringent for the personalized recall metric.

Desirable Surrogate Loss

- **Lesson learned:** Surrogate loss function should be defined in a **personalized** manner.


- **For each user**, we want the scores of positive items to be higher than those of the negative items
- We do *not* care about the score ordering across users.

- **Bayesian Personalized Ranking (BPR)** loss achieves



Loss Function: BPR Loss (1)

- **Bayesian Personalized Ranking (BPR) loss** is a personalized surrogate loss that aligns better with the recall@K metric.
- For each user $u^* \in U$, define the **rooted positive/negative edges** as
 - Positive edges rooted at u^*
 - $E(u^*) \equiv \{(u^*, v) \mid (u^*, v) \in E\}$
 - Negative edges rooted at u^*
 - $E_{\text{neg}}(u^*) \equiv \{(u^*, v) \mid (u^*, v) \in E_{\text{neg}}\}$



Note: The term “Bayesian” is not essential to the loss definition. The original paper [Rendle et al. 2009] considers the Bayesian prior over parameters (essentially acts as a parameter regularization), which we omit here.

Loss Function: BPR Loss (2)

- **Training objective:** For each user u^* , we want the scores of rooted positive edges $E(u^*)$ to be higher than those of rooted negative edges $E_{\text{neg}}(u^*)$.
- Aligns with the personalized nature of the recall metric.
- **BPR Loss for user u^* :**

Encouraged to be positive for each user

=positive edge score is higher than negative edge score

$$\text{Loss}(u^*) = \frac{1}{|E(u^*)| \cdot |E_{\text{neg}}(u^*)|} \underbrace{\sum_{(u^*, v_{\text{pos}}) \in E(u^*)} \sum_{(u^*, v_{\text{neg}}) \in E_{\text{neg}}(u^*)} -\log \left(\sigma \left(\overbrace{f_{\theta}(u^*, v_{\text{pos}}) - f_{\theta}(u^*, v_{\text{neg}})} \right) \right)}_{\text{Can be approximated using a mini-batch}}$$

Can be approximated using a mini-batch

- Final BPR Loss: $\frac{1}{|U|} \sum_{u^* \in U} \text{Loss}(u^*)$


Loss Function: BPR Loss (3)

■ Mini-batch training for the BPR loss:

- In each mini-batch, we sample a subset of users $U_{\text{mini}} \subset U$.
 - For each user $u^* \in U_{\text{mini}}$, we sample one positive item v_{pos} and a set of sampled negative items $V_{\text{neg}} = \{v_{\text{neg}}\}$.
- The mini-batch loss is computed as

$$\boxed{\frac{1}{|U_{\text{mini}}|} \sum_{u^* \in U_{\text{mini}}} \frac{1}{|V_{\text{neg}}|} \sum_{v_{\text{neg}} \in V_{\text{neg}}} -\log\left(\sigma\left(f_{\theta}(u^*, v_{\text{pos}}) - f_{\theta}(u^*, v_{\text{neg}})\right)\right)}$$

Average over users
in the mini-batch




Summary So Far

- **We have introduced**
 - Recall@ K as a metric for personalized recommendation
 - Embedding-based models
 - Three kinds of parameters to learn
 - **user encoder** to generate user embeddings
 - **item encoder** to generate item embeddings
 - **score function** to predict the user-item interaction likelihood.
 - Surrogate loss functions to achieve the high recall metric.
- Embedding-based models have achieved SoTA in recommender systems.
 - **Why do they work so well?**

Why Embedding Models Work?

- **Underlying idea:**
Collaborative filtering
 - Recommend items for a user by **collecting preferences of many other similar users.**
 - **Similar users tend to prefer similar items.**
- **Key question:** **How to capture similarity between users/items?**



Why Embedding Models Work?

- Embedding-based models can capture similarity of users/items!
 - **Low-dimensional embeddings *cannot* simply memorize all user-item interaction data.**
 - Embeddings are forced to **capture similarity between users/items to fit the data.**
 - This allows the models to make effective prediction on *unseen* user-item interactions.

This Lecture: GNNs for Recsys

- In this lecture, we teach two representative GNN approaches for recommender systems.
- **(1) Neural Graph Collab. Filtering (NGCF)** [Wang et al. 2019]
- **(2) LightGCN** [He et al. 2020]
 - Improve the conventional collaborative filtering models (i.e., shallow encoders) by explicitly modeling graph structure using GNNs.
 - Assumes no user/item features.
- **PinSAGE** [Ying et al. 2018]
 - Use GNNs to generate high-quality embeddings by simultaneously capturing rich node attributes (e.g., images) and the graph structure.


Stanford CS224W: Neural Graph Collaborative Filtering

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



Conventional Collaborative Filtering

- Conventional collaborative filtering model is based on **shallow encoders**:
 - No user/item features.
 - Use shallow encoders for users and items:
 - For every $u \in U$ and $v \in V$, we prepare shallow learnable embeddings $\mathbf{u}, \mathbf{v} \in \mathbb{R}^D$.
 - Score function for user u and item v is $f_{\theta}(\mathbf{u}, \mathbf{v}) \equiv \mathbf{z}_u^T \mathbf{z}_v$.



Limitations of Shallow Encoders

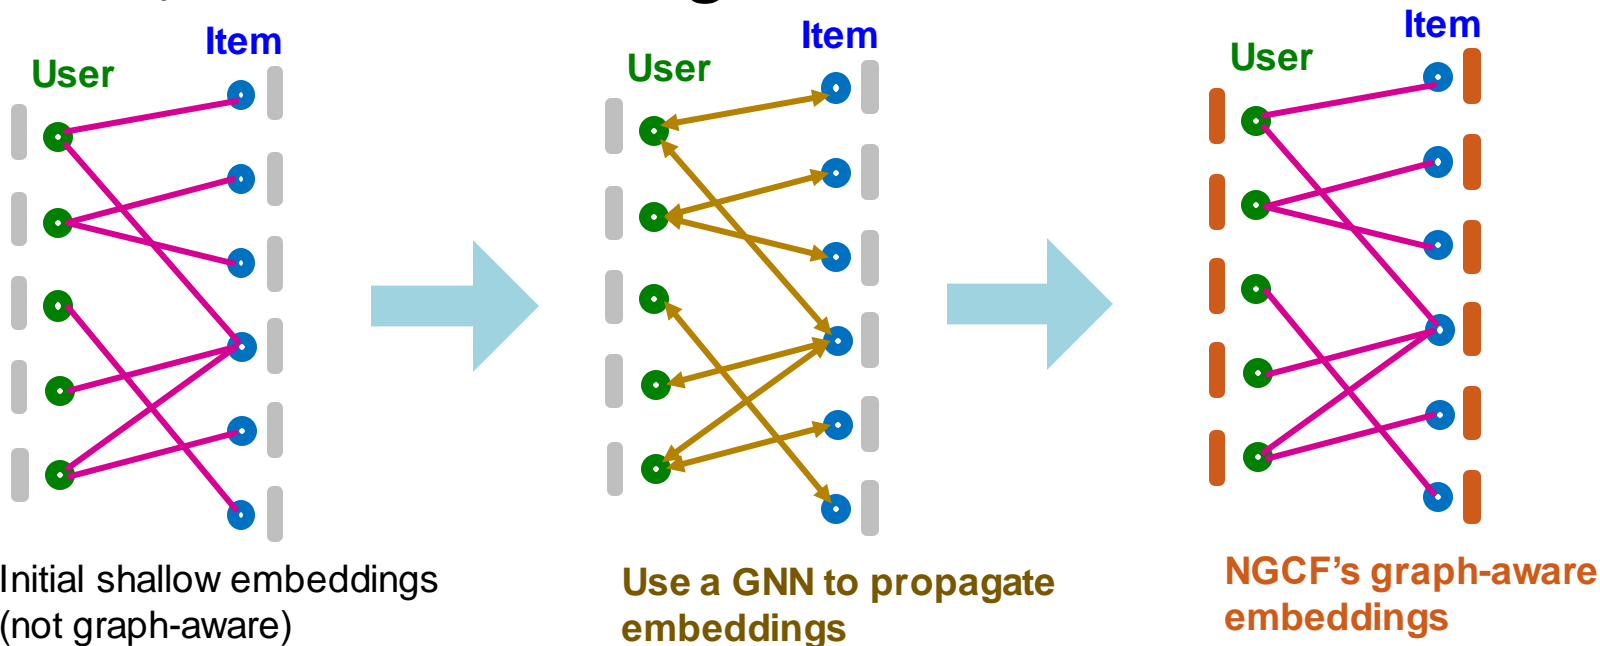
- The model itself does *not explicitly* capture graph structure
 - The graph structure is *only implicitly* captured in the training objective.
- Only the **first-order graph structure** (i.e., edges) is captured in the training objective.
 - **High-order graph structure** (e.g., K -hop paths between two nodes) is *not explicitly captured*.

Motivation

- We want a model that...
 - **explicitly captures graph structure** (beyond implicitly through the training objective)
 - captures **high-order graph structure** (beyond the first-order edge connectivity structure)
- **GNNs are a natural approach to achieve both!**
 - **Neural Graph Collaborative Filtering (NGCF)** [Wang et al. 2019]
 - **LightGCN** [He et al. 2020]
 - A simplified and improved version of NGCF


NGCF: Overview

- **Neural Graph Collaborative Filtering (NGCF)** *explicitly* incorporates high-order graph structure when generating user/item embeddings.
- **Key idea:** Use a GNN to generate graph-aware user/item embeddings.




NGCF Framework

- **Given:** User-item bipartite graph.
- **NGCF framework:**
 - Prepare shallow learnable embedding for each node.
 - Use multi-layer GNNs to propagate embeddings along the bipartite graph.
 - High-order graph structure is captured.
 - Final embeddings are *explicitly* graph-aware!
- **Two kinds of learnable params are jointly learned:**
 - Shallow user/item embeddings
 - GNN's parameters



Initial Node Embeddings

- Set the shallow learnable embeddings as the initial node features:
 - For every user $u \in U$, set $\mathbf{h}_u^{(0)}$ as the user's shallow embedding.
 - For every item $v \in V$, set $\mathbf{h}_v^{(0)}$ as the item's shallow embedding.



Neighbor Aggregation

- Iteratively update node embeddings using neighboring embeddings.


$$\mathbf{h}_v^{(k+1)} = \text{COMBINE}\left(\mathbf{h}_v^{(k)}, \text{AGGR}\left(\left\{\mathbf{h}_u^{(k)}\right\}_{u \in N(v)}\right)\right)$$

$$\mathbf{h}_u^{(k+1)} = \text{COMBINE}\left(\mathbf{h}_u^{(k)}, \text{AGGR}\left(\left\{\mathbf{h}_v^{(k)}\right\}_{v \in N(u)}\right)\right)$$

High-order graph structure is captured through iterative neighbor aggregation.

Different architecture choices are possible for AGGR and COMBINE.

- AGGR(\cdot) can be MEAN(\cdot)
- COMBINE(\mathbf{x}, \mathbf{y}) can be $\text{ReLU}(\text{Linear}(\text{Concat}(\mathbf{x}, \mathbf{y})))$




Final Embeddings and Score Function

- After K rounds of neighbor aggregation, we get the **final user/item embeddings** $\mathbf{h}_u^{(K)}$ and $\mathbf{h}_v^{(K)}$.

- For all $u \in U$, $v \in V$, we set
 $\mathbf{u} \leftarrow \mathbf{h}_u^{(K)}$, $\mathbf{v} \leftarrow \mathbf{h}_v^{(K)}$.

- Score function is the inner product

$$\text{score}(u, v) = \mathbf{u}^T \mathbf{v}$$



NGCF: Summary

- Conventional collaborative filtering uses shallow user/item embeddings.
 - The embeddings do *not explicitly model graph structure*.
 - The training objective *does not model high-order graph structure*.
- NGCF uses a GNN to propagate the shallow embeddings.
 - The embeddings are *explicitly aware of high-order graph structure*.

Stanford CS224W: LightGCN

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



LightGCN: Motivation (1)

- **Recall:** NGCF jointly learns two kinds of parameters:
 - Shallow user/item embeddings
 - GNN's parameters
- **Observation:** Shallow learnable embeddings are already quite expressive.
 - They are learned for every (user/item) node.
 - Most of the parameter counts are in shallow embeddings when N (#nodes) $\gg D$ (embedding dimensionality)
 - Shallow embeddings: $O(ND)$.
 - GNN: $O(D^2)$.
 - The GNN parameters may not be so essential for performance.


LightGCN: Motivation (2)

- Can we simplify the GNN used in NGCF (e.g., remove its learnable parameters)?
 - **Answer:** Yes!
 - **Bonus:** Simplification improves the recommendation performance!
- **Overview of the idea:**
 - Adjacency matrix for a bipartite graph
 - Matrix formulation of GCN
 - Simplification of GCN by removing non-linearity
 - Related: SGC for scalable GNN [Wu et al. 2019]

Adjacency and Embedding Matrices



- **Adjacency matrix** of a (undirected) bipartite graph.
- **Shallow embedding matrix**.

Adjacency matrix A



$R_{uv} = 1$ if
user u
interacts
with item v ,
 $R_{uv} = 0$
otherwise.

Embedding
matrix E



Matrix Formulation of GCN

- **Recall**: The diffusion matrix of C&S.
- Let \mathbf{D} be the degree matrix of \mathbf{A} .
- Define the normalized adjacency matrix $\tilde{\mathbf{A}}$ as

$$\tilde{\mathbf{A}} \equiv \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$$


- Let $\mathbf{E}^{(k)}$ be the embedding matrix at k -th layer.
- Each layer of GCN's aggregation can be written in a matrix form:

$$\mathbf{E}^{(k+1)} = \text{ReLU}(\tilde{\mathbf{A}} \mathbf{E}^{(k)} \mathbf{W}^{(k)})$$

Neighbor aggregation Learnable linear transformation

Note: Different from the original GCN, self-connection is omitted here.

Matrix of node embeddings $\mathbf{E}^{(k)}$



Each row stores node embedding

Simplifying GCN (1)

- Simplify GCN by removing ReLU non-linearity:

$$E^{(k+1)} = \tilde{A} E^{(k)} W^{(k)} \quad \text{Original idea from SGC [Wu et al. 2019]}$$

- The final node embedding matrix is given as

$$\begin{aligned} E^{(K)} &= \tilde{A} \underbrace{E^{(K-1)}}_{\text{pink}} W^{(K-1)} \\ &= \tilde{A} \left(\underbrace{\tilde{A} E^{(K-2)}}_{\text{green}} W^{(K-2)} \right) W^{(K-1)} \\ &= \tilde{A} \left(\tilde{A} \left(\dots \left(\underbrace{\tilde{A} E^{(0)}}_{\text{green}} W^{(0)} \right) \dots \right) W^{(K-2)} \right) W^{(K-1)} \\ &= \tilde{A}^K \underbrace{E}_{\text{blue}} \left(\underbrace{W^{(0)} \dots W^{(K-1)}}_{\text{orange}} \right) \end{aligned}$$

Set E as input embedding $E^{(0)}$

Simplifying GCN (2)

- Removing ReLU significantly simplifies GCN!

$$\mathbf{E}^{(K)} = \boxed{\tilde{\mathbf{A}}^K \mathbf{E}} \mathbf{W} \quad \mathbf{W} \equiv \mathbf{W}^{(0)} \dots \mathbf{W}^{(K-1)}$$

Diffusing node embeddings
along the graph

(similar to C&S that diffuses soft
labels along the graph)

- **Algorithm:** Apply $\mathbf{E} \leftarrow \tilde{\mathbf{A}} \mathbf{E}$ for K times.
 - Each matrix multiplication diffuses the current embeddings to their one-hop neighbors.
 - **Note:** $\tilde{\mathbf{A}}^K$ is dense and never gets materialized. Instead, the above iterative matrix-vector product is used to compute $\tilde{\mathbf{A}}^K \mathbf{E}$.

Multi-Scale Diffusion

- We can consider **multi-scale diffusion**

$$\alpha_0 E^{(0)} + \alpha_1 E^{(1)} + \alpha_2 E^{(2)} + \dots + \alpha_K E^{(K)}$$


- The above includes embeddings diffused at multiple hop scales.
- $\alpha_0 E^{(0)} = \alpha_0 \tilde{A}^0 E^{(0)}$ acts as a self-connection (that is omitted in the definition \tilde{A})
- The coefficients, $\alpha_0, \dots, \alpha_K$, are hyper-parameters.
- For simplicity, LightGCN uses the uniform coefficient, i.e., $\alpha_k = \frac{1}{K+1}$ for $k = 0, \dots, K$.

LightGCN: Model Overview (2)


Iteratively diffuse embedding matrix E using \tilde{A}

For $k = 0 \dots K - 1$,


Embedding
matrix $E^{(k+1)}$



Normalized Adj. matrix \tilde{A}
(self-loop omitted)



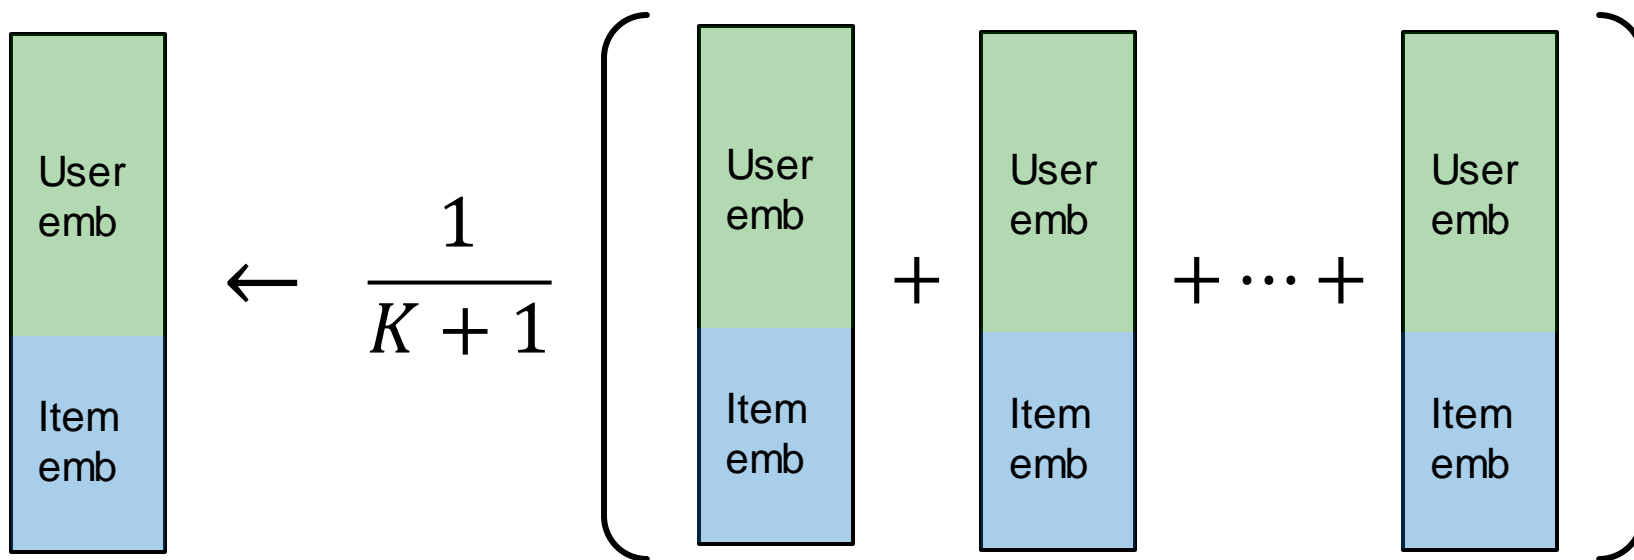
Embedding
matrix $E^{(k)}$
($E^{(0)}$ is set to E)



LightGCN: Model Overview (3)

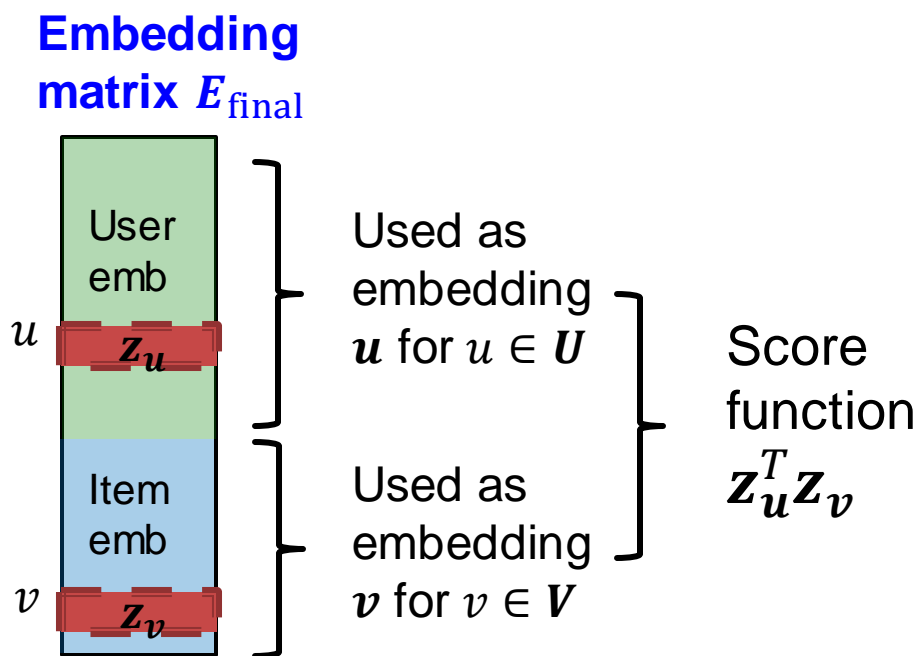
- Average the embedding matrices at different scales.

Embedding
matrix E_{final}



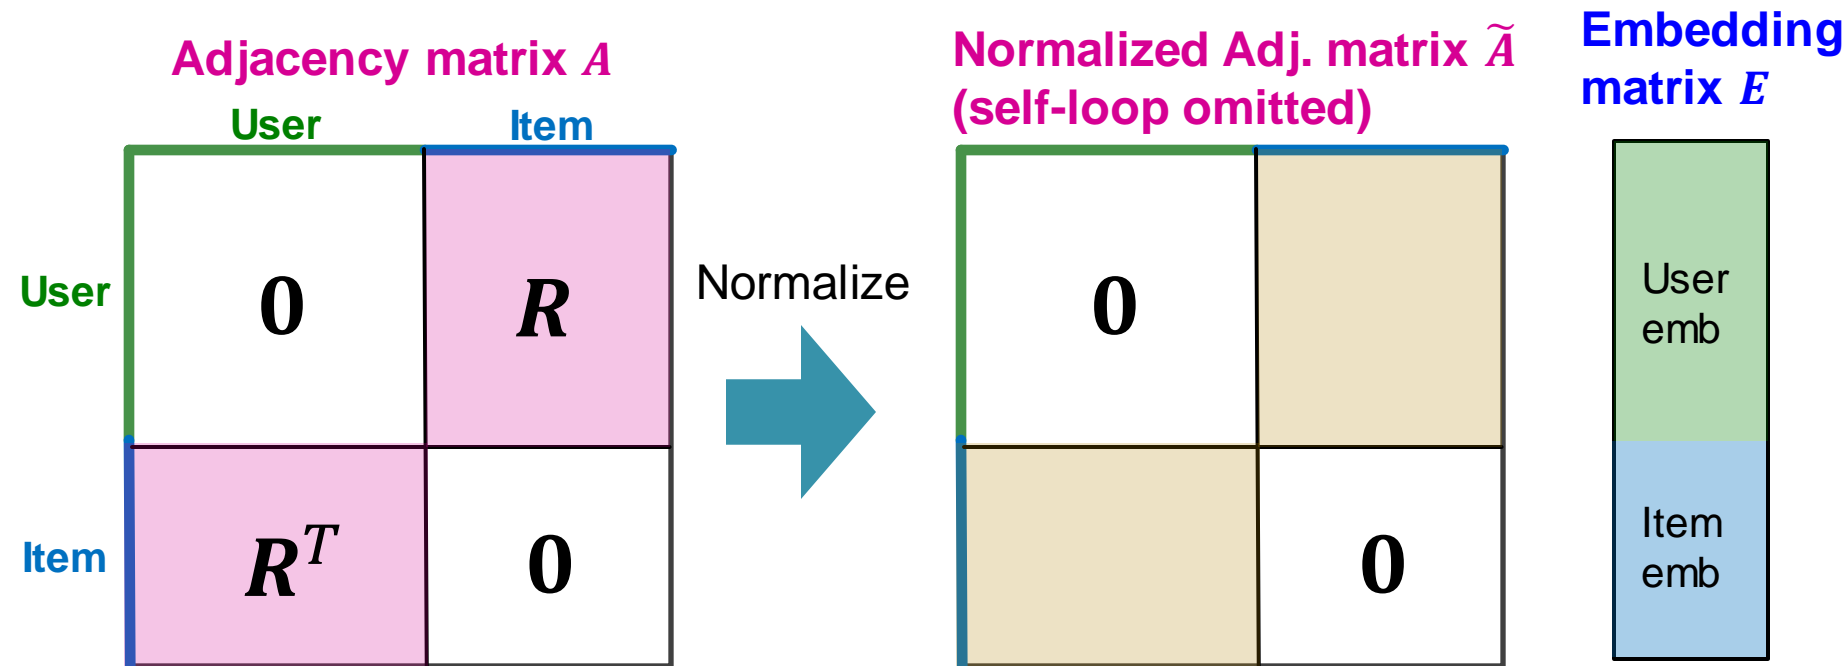
LightGCN: Model Overview (4)

- **Score function:**
 - Use user/item vectors from E_{final} to score user-item interaction



LightGCN: Model Overview (1)

- **Given:**
 - Adjacency matrix A
 - Initial learnable embedding matrix E



LightGCN: Intuition

- **Question:** Why does the simple diffusion propagation work well?
- **Answer:** The diffusion directly encourages the embeddings of similar users/items to be similar.
 - Similar users share many common neighbors (items) and are expected to have similar future preferences (interact with similar items).

LightGCN and GCN/C&S

- The embedding propagation of LightGCN is closely related to GCN/C&S.
- **Recall:** GCN/C&S (neighbor aggregation part)

$$\mathbf{h}_v^{(k+1)} = \sum_{u \in N(v)} \frac{1}{\sqrt{d_u} \sqrt{d_v}} \cdot \mathbf{h}_u^{(k)}$$

Node degree

- Self-loop is added in the neighborhood definition.
- LightGCN uses the same equation except that
 - Self-loop is *not* added in the neighborhood definition.
 - Final embedding takes the average of embeddings from all the layers: $\mathbf{h}_v = \frac{1}{K+1} \sum_{k=0}^K \mathbf{h}_v^{(k)}$.

LightGCN and MF: Comparison

- Both LightGCN and shallow encoders **learn a unique embedding for each user/item.**
- The difference is that LightGCN uses the *diffused* user/item embeddings for scoring.
- LightGCN performs better than shallow encoders but are also more computationally expensive due to the additional diffusion step.
 - The final embedding of a user/item is obtained by aggregating embeddings of its multi-hop neighboring nodes.

LightGCN: Summary

- LightGCN simplifies NGCF by **removing the learnable parameters of GNNs.**
- **Learnable parameters are all in the shallow input node embeddings.**
 - Diffusion propagation only involves matrix-vector multiplication.
 - The simplification leads to better empirical performance than NGCF.

Stanford CS224W: PinSAGE

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



PinSAGE: Pin Embedding


- ❑ Unifies visual, textual, and graph information.
- ❑ The largest industry deployment of a Graph Convolutional Networks.
- ❑ Huge Adoption across Pinterest
- ❑ Works for fresh content and is available in a few seconds after pin creation



Application: Pinterest

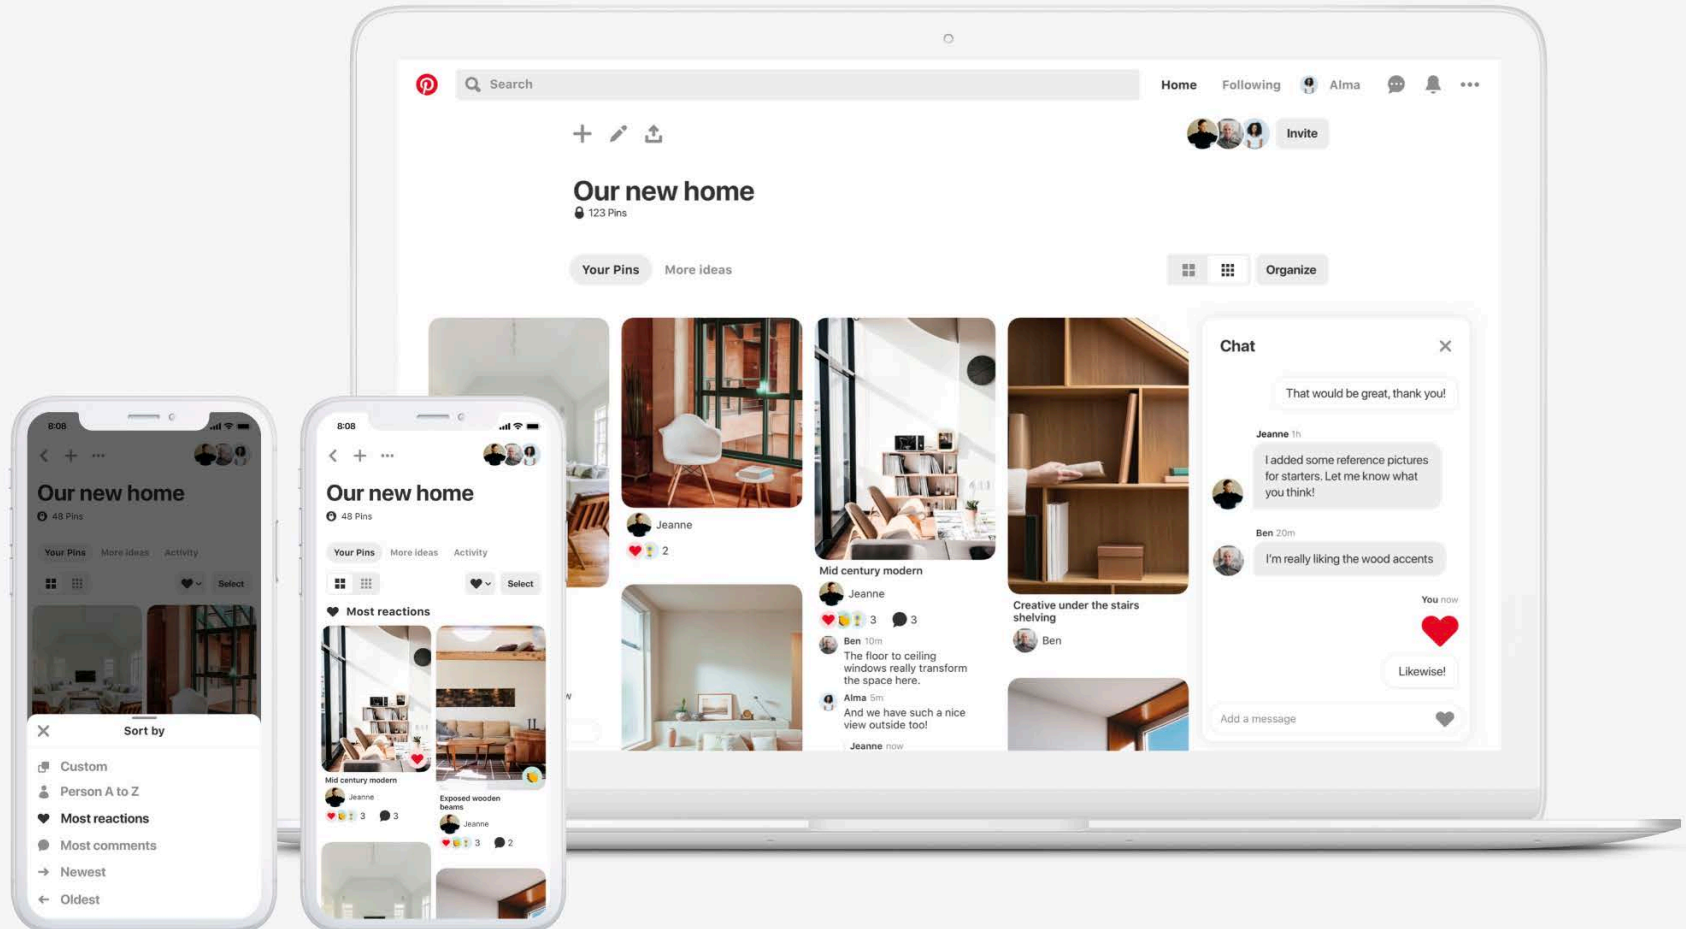
PinSage graph convolutional network:

- **Goal:** Generate embeddings for nodes in a large-scale Pinterest graph containing billions of objects
- **Key Idea:** Borrow information from nearby nodes
 - E.g., bed rail Pin might look like a garden fence, but gates and beds are rarely adjacent in the graph



- Pin embeddings are essential to various tasks like recommendation of Pins, classification, ranking
 - Services like “Related Pins”, “Search”, “Shopping”, “Ads”


Harnessing Pins and Boards



source: Pinterest newsroom

PinSAGE: Graph Neural Network


- Graph has tens of billions of nodes and edges
- Further resolves embeddings across the Pinterest graph



PinSAGE: Methods for Scaling Up

- In addition to the GNN model, the PinSAGE paper introduces several methods to scale the GNN to a billion-scale recommender system (e.g., Pinterest).
 - Shared negative samples across users in a mini-batch
 - Hard negative samples
 - Curriculum learning
 - Mini-batch training of GNNs on a large-graph (to be covered in the future lecture)

PinSAGE Model



Task: Recommend related pins to users

Learn node embeddings z_i such that

$$d(z_{cake1}, z_{cake2}) < d(z_{cake1}, z_{sweater})$$




Source pin



SUCCESSFUL
RECOMMENDATION



BAD RECOMMENDATION



Training Data

1+B repin pairs:


- From Related Pins surface
- Capture semantic relatedness
- Goal: Embed such pairs to be “neighbors”

Example positive training pairs (Q,X):



Shared Negative Samples (1)

- **Recall:** In BPR loss, for each user $u^* \in \mathbf{U}_{\text{mini}}$, we sample one positive item v_{pos} and a set of sampled negative items $\mathbf{V}_{\text{neg}} = \{v_{\text{neg}}\}$.
- Using more negative samples per user improves the recommendation performance, but is also expensive.
 - We need to generate $|\mathbf{U}_{\text{mini}}| \cdot |\mathbf{V}_{\text{neg}}|$ embeddings for negative nodes.
 - We need to apply $|\mathbf{U}_{\text{mini}}| \cdot |\mathbf{V}_{\text{neg}}|$ GNN computational graphs (see right), which is expensive.



Shared Negative Samples (2)

- **Key idea:** We can share the same set of negative samples $V_{\text{neg}} = \{v_{\text{neg}}\}$ *across all users* U_{mini} in the mini-batch.
- This way, we only need to generate $|V_{\text{neg}}|$ embeddings for negative nodes.
 - This saves the node embedding generation computation **by a factor of $|U_{\text{mini}}|$** !
 - Empirically, the performance stays similar to the non-shared negative sampling scheme.

Hard Negatives (1)

- **Challenge**: Industrial recsys needs to make **extremely fine-grained predictions**.
 - #Total items: Up to billions.
 - #Items to recommend for each user: 10 to 100.
- **Issue**: The shared negative items are randomly sampled from all items
 - Most of them are “**easy negatives**”, i.e., a model does not need to be fine-grained to distinguish them from positive items.
- We need a way to sample “**hard negatives**” to force the model to be fine-grained!

PinSAGE: Curriculum Learning

- **Idea:** use harder and harder negative samples
- Include more and more hard negative samples for each epoch



Source pin



Positive



Easy negative



Hard negative

Curriculum Learning

- **Key insight:** It is effective **to make the negative samples *gradually harder* in the process of training.**
- At n -th epoch, we add $n - 1$ hard negative items.
 - #(Hard negatives) gradually increases in the process of training.
- The model will gradually learn to make finer-grained predictions.



Hard Negatives (2)

- For each user node, the **hard negatives** are item nodes that are close (but not connected) to the user node in the graph.
- Hard negatives for user $u \in U$ are obtained as follows:
 - Compute personalized page rank (PPR) for user u .
 - Sort items in the descending order of their PPR scores.
 - Randomly sample item nodes that are ranked high but not too high, e.g., 2000th — 5000th.
 - Item nodes that are close but not too close (connected) to the user node.
- The hard negatives for each user are used in addition to the shared negatives.

PinSAGE: Negative Sampling

(q, p) positive pairs are given but various methods to sample negatives to form (q, p, n)

- Distance Weighted Sampling ([Wu et al., 2017](#))
 - Sample negatives so that query-negative distance distribution is approx $U[0.5, 1.4]$



(b) Sample distribution for different strategies.

PinSAGE: Summary

- **PinSAGE uses GNNs** to generate high-quality user/item embeddings that **capture both the rich node attributes and graph structure**.
- The PinSAGE model is effectively trained using sophisticated **negative sampling strategies**.
- PinSAGE is **successfully deployed at Pinterest**, a billion-scale image content recommendation service.
 - **Uncovered in this lecture**: How to **scale up GNNs to large-scale graphs**. Will be covered in a later lecture.